

The following (fixed-size) array type exists:

- `<type>[M]` : a fixed-length array of the given fixed-length type.

The following non-fixed-size types exist:

- `bytes` : dynamic sized byte sequence.
- `string` : dynamic sized unicode string assumed to be UTF-8 encoded.
- `<type>[]` : a variable-length array of the given fixed-length type.

Formal Specification of the Encoding

We will now formally specify the encoding, such that it will have the following properties, which are especially useful if some arguments are nested arrays:

Properties:

- The number of reads necessary to access a value is at most the depth of the value inside the argument array structure, i.e. four reads are needed to retrieve `a_i[k][l][r]` . In a previous version of the ABI, the number of reads scaled linearly with the total number of dynamic parameters in the worst case.
- The data of a variable or array element is not interleaved with other data and it is relocatable, i.e. it only uses relative "addresses"

We distinguish static and dynamic types. Static types are encoded in-place and dynamic types are encoded at a separately allocated location after the current block.

Definition: The following types are called "dynamic":

- `bytes`
- `string`
- `T[]` for any `T`
- `T[k]` for any dynamic `T` and any `k > 0`

All other types are called "static".

Definition: `len(a)` is the number of bytes in a binary string `a` . The type of `len(a)` is assumed to be `uint256` .

We define `enc` , the actual encoding, as a mapping of values of the ABI types to binary strings such that `len(enc(X))` depends on the value of `X` if and only if the type of `X` is dynamic.

Definition: For any ABI value `X` , we recursively define `enc(X)` , depending on the type of `X` being

- `T[k]` for any `T` and `k` :

$$\text{enc}(X) = \text{head}(X[0]) \dots \text{head}(X[k-1]) \text{tail}(X[0]) \dots \text{tail}(X[k-1])$$

where `head` and `tail` are defined for `X[i]` being of a static type as `head(X[i]) = enc(X[i])` and `tail(X[i]) = ""` (the empty string) and as `head(X[i]) = enc(len(head(X[0]) ... head(X[k-1]) tail(X[0]) ... tail(X[i-1])))` `tail(X[i]) = enc(X[i])` otherwise.

Note that in the dynamic case, `head(X[i])` is well-defined since the lengths of the head parts only depend on the types and not the values. Its value is the offset of the beginning of `tail(X[i])` relative to the start of `enc(X)` .

- `T[]` where `X` has `k` elements (`k` is assumed to be of type `uint256`):

$$\text{enc}(X) = \text{enc}(k) \text{enc}([X[1], \dots, X[k]])$$

i.e. it is encoded as if it were an array of static size `k` , prefixed with the number of elements.

DEV Technologies

- [RLP Encoding](#)
- [RLPx Node Discovery Protocol](#)
- [EIP2p Wire Protocol](#)
- [EIP2p Whitepaper \(WiP\)](#)
- [Web3 Secret Storage](#)

Ethereum Technologies

- [Patricia Tree](#)
- [Wire protocol](#)
- [Light client protocol](#)
- [Subtleties](#)
- [Solidity, Docs & ABI](#)
- [NatSpec Format](#)
- [Contract ABI](#)
- [Bad Block Reporting](#)
- [Bad Chain Canary](#)
- [Extra Data](#)
- [Brain Wallet](#)

Ethash/Dashimoto

- [Ethash](#)
- [Ethash C API](#)
- [Ethash DAG](#)

Whisper

- [Whisper Proposal](#)
- [Whisper Overview](#)
- [PoC-1 Wire protocol](#)
- [PoC-2 Wire protocol](#)
- [PoC-2 Whitepaper](#)

Misc

- [Hard Problems of Cryptocurrency](#)
- [Chain Fibers](#)
- [Glossary](#)

Clone this wiki locally

<https://github.com/ethereum/wiki>



Clone in Desktop

[illegible]

- In total:

[illegible][illegible]

If we wanted to call `sam` with the arguments `"dave"`, `true` and `[1,2,3]`, we would pass 292 bytes total, broken down into:

- In total:

[illegible]

Use of Dynamic Types

A call to a function with the signature `f(uint,uint32[],bytes10,bytes)` with values `(0x123, [0x456, 0x789], "1234567890", "Hello, world!")` is encoded in the following way:

We take the first four bytes of `sha3("f(uint256,uint32[],bytes10,bytes)")` , i.e. `0x8be65246` . Then we encode the head parts of all four arguments. For the static types `uint256` and `bytes10` , these are directly the values we want to pass, whereas for the dynamic types `uint32[]` and `bytes` , we use the offset in bytes to the start of their data area, measured from the start of the value encoding (i.e. not counting the first four bytes containing the hash of the function signature). These are:

- [illegible]

After this, the data part of the first dynamic argument, `[0x456, 0x789]` follows:

- [illegible]

Finally, we encode the data part of the second dynamic argument, "Hello, world!" :

- `0x00d` (number of elements (bytes in this case): 13)
- `0x48656c6c6f2c20776f7226c6421000000000000000000000000000000000000` ("Hello, world!" padded to 32 bytes on the right)

All together, the encoding is (newline after function selector and each 32-bytes for clarity):

[illegible]

Events

Events are an abstraction of the Ethereum logging/event-watching protocol. Log entries provide the contract's address, a series of up to four topics and some arbitrary length binary data. Events leverage the existing function ABI in order to interpret this (together with an interface spec) as a properly typed structure.

Given an event name and series of event parameters, we split them into two sub-series: those which are indexed and those which are not. Those which are indexed, which may number up to

3, are used alongside the Keccak hash of the event signature to form the topics of the log entry. Those which as not indexed form the byte array of the event.

In effect, a log entry using this ABI is described as:

- `address` : the address of the contract (intrinsically provided by Ethereum);
- `topics[0]` : `keccak(EVENT_NAME+"("+EVENT_ARGS.map(canonical_type_of).join(",")+"))` (`canonical_type_of` is a function that simply returns the canonical type of a given argument, e.g. for `uint indexed foo` , it would return `uint256`). If the event is declared as `anonymous` the `topics[0]` is not generated;
- `topics[n]` : `EVENT_INDEXED_ARGS[n - 1]` (`EVENT_INDEXED_ARGS` is the series of `EVENT_ARGS` that are indexed);
- `data` : `abi_serialise(EVENT_NON_INDEXED_ARGS)` (`EVENT_NON_INDEXED_ARGS` is the series of `EVENT_ARGS` that are not indexed, `abi_serialise` is the ABI serialisation function used for returning a series of typed values from a function, as described above).

JSON

The JSON format for a contract's interface is given by an array of function and/or event descriptions. A function description is a JSON object with the fields:

- `type` : "function" or "constructor" (can be omitted, defaulting to function);
- `name` : the name of the function (only present for function types);
- `inputs` : an array of objects, each of which contains:
 - `name` : the name of the parameter;
 - `type` : the canonical type of the parameter.
- `outputs` : an array of objects similar to `inputs` , can be omitted.

An event description is a JSON object with fairly similar fields:

- `type` : always "event"
- `name` : the name of the event;
- `inputs` : an array of objects, each of which contains:
 - `name` : the name of the parameter;
 - `type` : the canonical type of the parameter.
 - `indexed` : `true` if the field is part of the log's topics, `false` if it one of the log's data segment.
- `anonymous` : `true` if the event was declared as `anonymous` .

For example,

```
contract Test {
function Test(){ b = 0x12345678901234567890123456789012; }
event Event(uint indexed a, bytes32 b)
event Event2(uint indexed a, bytes32 b)
function foo(uint a) { Event(a, b); }
bytes32 b;
}
```

would result in the JSON:

```
[{
  "type":"event",
  "inputs": [{ "name":"a","type":"uint256","indexed":true},{ "name":"b","type":"bytes32" }],
  "name":"Event"
}, {
  "type":"event",
  "inputs": [{ "name":"a","type":"uint256","indexed":true},{ "name":"b","type":"bytes32" }],
  "name":"Event2"
}]
```

```
}, {
  "type": "event",
  "inputs": [{ "name": "a", "type": "uint256", "indexed": true }, { "name": "b", "type": "bytes32" }],
  "name": "Event2"
}, {
  "type": "function",
  "inputs": [{ "name": "a", "type": "uint256" }],
  "name": "foo",
  "outputs": []
}]
```

Example Javascript Usage

```
var Test = eth.contract([
  {
    "type": "event",
    "inputs": [{ "name": "a", "type": "uint256", "indexed": true }, { "name": "b", "type": "bytes32" }],
    "name": "Event"
  }, {
    "type": "event",
    "inputs": [{ "name": "a", "type": "uint256", "indexed": true }, { "name": "b", "type": "bytes32" }],
    "name": "Event2"
  }, {
    "type": "function",
    "inputs": [{ "name": "a", "type": "uint256" }],
    "name": "foo",
    "outputs": []
  }
]);
var theTest = new Test(addrTest);

// examples of usage:
// every log entry ("event") coming from theTest (i.e. Event & Event2):
var f0 = eth.filter(theTest);
// just log entries ("events") of type "Event" coming from theTest:
var f1 = eth.filter(theTest.Event);
// also written as
var f1 = theTest.Event();
// just log entries ("events") of type "Event" and "Event2" coming from theTest:
var f2 = eth.filter([theTest.Event, theTest.Event2]);
// just log entries ("events") of type "Event" coming from theTest with indexed parameter 'a' equal to 69:
var f3 = eth.filter(theTest.Event, { 'a': 69 });
// also written as
var f3 = theTest.Event({ 'a': 69 });
// just log entries ("events") of type "Event" coming from theTest with indexed parameters 'a' and 'b' equal to 69 and 42:
var f4 = eth.filter(theTest.Event, { 'a': [69, 42] });
// also written as
var f4 = theTest.Event({ 'a': [69, 42] });

// options may also be supplied as a second parameter with `earliest`, `latest`, and `fromBlock`
var options = { 'max': 100 };
var f4 = theTest.Event({ 'a': [69, 42] }, options);

var trigger;
f4.watch(trigger);

// call foo to make an Event:
theTest.foo(69);

// would call trigger like:
// trigger(theTest.Event, { 'a': 69, 'b': '0x12345678901234567890123456789012' }, n)
// where n is the block number that the event triggered in.
```

Implementation:

```
// e.g. f4 would be similar to:
web3.eth.filter({ 'max': 100, 'address': theTest.address, 'topics': [ [69, 42] ] })
```

```
// except that the resultant data would need to be converted from the basic log e
{
  'address': theTest.address,
  'topics': [web3.sha3("Event(uint256,bytes32)", 0x00...0045 /* 69 in hex format
  'data': '0x12345678901234567890123456789012',
  'number': n
}
// into data good for the trigger, specifically the three fields:
Test.Event // derivable from the first topic
{'a': 69, 'b': '0x12345678901234567890123456789012'} // derivable from the 'ind
n // from the 'number'
```

Event result:

```
[ {
  'event': Test.Event,
  'args': {'a': 69, 'b': '0x12345678901234567890123456789012'},
  'number': n
},
{ ...
} ...
]
```

JUST DONE! [develop branch]

NOTE: THIS IS OLD - IGNORE IT unless reading for historical purposes

- Internal LogFilter, log-entry matching mechanism and eth_installFilter needs to support matching multiple values (OR semantics) *per* topic index (at present it will only match topics with AND semantics and set-inclusion, not per-index).

i.e. at present you can only ask for each of a number of given topic values to be matched throughout each topic:

- `topics: [69, 42, "Gav"]` would match against logs with 3 topics `[42, 69, "Gav"]` , `["Gav", 69, 42]` but **not** against logs with topics `[42, 70, "Gav"]` .

we need to be able to provide one of a number of topic values, and, each of these options for each topic index:

- `topics: [[69, 42], [] /* anything */, "Gav"]` should match against logs with 3 topics `[42, 69, "Gav"]` , `[42, 70, "Gav"]` but **not** against `["Gav", 69, 42]` .

[\[English | Deutsch | Español | Français | 日本語 | Română | فارسی | Italiano | 한국어 | 中文 \]](#)